

Navigation

In this notebook, you will learn how to use the Unity ML-Agents environment for the first project of the [Deep Reinforcement Learning Nanodegree](#).

1. Start the Environment

We begin by importing some necessary packages. If the code cell below returns an error, please revisit the project instructions to double-check that you have installed [Unity ML-Agents](#) and [NumPy](#).

In [1]:

```
from unityagents import UnityEnvironment
import numpy as np
from dqn_agent import Agent

import matplotlib.pyplot as plt
%matplotlib inline
```

Next, we will start the environment! **Before running the code cell below**, change the `file_name` parameter to match the location of the Unity environment that you downloaded.

- **Mac:** "path/to/Banana.app"
- **Windows (x86):** "path/to/Banana_Windows_x86/Banana.exe"
- **Windows (x86_64):** "path/to/Banana_Windows_x86_64/Banana.exe"
- **Linux (x86):** "path/to/Banana_Linux/Banana.x86"
- **Linux (x86_64):** "path/to/Banana_Linux/Banana.x86_64"
- **Linux (x86, headless):** "path/to/Banana_Linux_NoVis/Banana.x86"
- **Linux (x86_64, headless):** "path/to/Banana_Linux_NoVis/Banana.x86_64"

For instance, if you are using a Mac, then you downloaded `Banana.app`. If this file is in the same folder as the notebook, then the line below should appear as follows:

```
env = UnityEnvironment(file_name="Banana.app")
```

In [2]:

```
env = UnityEnvironment(file_name="Banana.app")
```

```
INFO:unityagents:
'Academy' started successfully!
Unity Academy name: Academy
    Number of Brains: 1
    Number of External Brains : 1
    Lesson number : 0
    Reset Parameters :

Unity brain name: BananaBrain
    Number of Visual Observations (per agent): 0
    Vector Observation space type: continuous
    Vector Observation space size (per agent): 37
    Number of stacked Vector Observation: 1
    Vector Action space type: discrete
    Vector Action space size (per agent): 4
    Vector Action descriptions: , , ,
```

Environments contain **brains** which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

In [3]:

```
# get the default brain
brain_name = env.brain_names[0]
```

```
brain_name = env.brain_names[0]
brain = env.brains[brain_name]
```

2. Examine the State and Action Spaces

The simulation contains a single agent that navigates a large environment. At each time step, it has four actions at its disposal:

- 0 - walk forward
- 1 - walk backward
- 2 - turn left
- 3 - turn right

The state space has 37 dimensions and contains the agent's velocity, along with ray-based perception of objects around agent's forward direction. A reward of +1 is provided for collecting a yellow banana, and a reward of -1 is provided for collecting a blue banana.

Run the code cell below to print some information about the environment.

In [4]:

```
# reset the environment
env_info = env.reset(train_mode=True)[brain_name]

# number of agents in the environment
print('Number of agents:', len(env_info.agents))

# number of actions
action_size = brain.vector_action_space_size
print('Number of actions:', action_size)

# examine the state space
state = env_info.vector_observations[0]
print('States look like:', state)
state_size = len(state)
print('States have length:', state_size)
```

```
Number of agents: 1
Number of actions: 4
States look like: [1.          0.          0.          0.          0.84408134  0.
 0.          1.          0.          0.0748472  0.          1.
 0.          0.          0.25755    1.          0.          0.
 0.          0.74177343  0.          1.          0.          0.
 0.25854847  0.          0.          1.          0.          0.09355672
 0.          1.          0.          0.          0.31969345  0.
 0.          ]
States have length: 37
```

In [6]:

```
agent = Agent(state_size=state_size, action_size=action_size, seed=0)
agent.qnetwork_local.load_state_dict(torch.load('checkpoint.pth'))

def dqn(n_episodes=1500, max_t=1000, eps_start=1.0, eps_end=0.01, eps_decay=0.995):
    """Deep Q-Learning.

    Params
    =====
    n_episodes (int): maximum number of training episodes
    max_t (int): maximum number of timesteps per episode
    eps_start (float): starting value of epsilon, for epsilon-greedy action selection
    eps_end (float): minimum value of epsilon
    eps_decay (float): multiplicative factor (per episode) for decreasing epsilon
    """
    scores = [] # list containing scores from each episode
    scores_window = deque(maxlen=100) # last 100 scores
    eps = eps_start # initialize epsilon
    for i_episode in range(1, n_episodes+1):
        env_info = env.reset(train_mode=True)[brain_name] # reset the environment
        state = env_info.vector_observations[0]
        score = 0
        for t in range(max_t):
            action = agent.act(state, eps)
```

```

env_info = env.step(action)[brain_name]      # send the action to the environment
next_state = env_info.vector_observations[0]  # get the next state
reward = env_info.rewards[0]                # get the reward
done = env_info.local_done[0]
agent.step(state, action, reward, next_state, done)
score += reward                              # update the score
state = next_state

if done:
    break
scores_window.append(score)                  # save most recent score
scores.append(score)                        # save most recent score
eps = max(eps_end, eps_decay*eps)          # decrease epsilon
print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)), end=
""))

if i_episode % 100 == 0:
    print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))


if np.mean(scores_window) >= 15:
    print('\nEnvironment solved in {:d} episodes! \tAverage Score: {:.2f}'.format(i_episode-
100, np.mean(scores_window)))
    torch.save(agent.qnetwork_local.state_dict(), 'checkpoint.pth')
    break
return scores
scores = dqn()
# plot the scores
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores)
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.show()

```

```

Episode 100 Average Score: 4.47
Episode 200 Average Score: 11.78
Episode 300 Average Score: 14.16
Episode 400 Average Score: 15.84
Episode 500 Average Score: 16.71
Episode 600 Average Score: 16.02
Episode 700 Average Score: 16.76
Episode 800 Average Score: 17.66
Episode 900 Average Score: 16.28
Episode 1000 Average Score: 16.63
Episode 1100 Average Score: 17.02
Episode 1200 Average Score: 16.99
Episode 1300 Average Score: 16.05
Episode 1400 Average Score: 16.66
Episode 1500 Average Score: 16.09

```



When finished, you can close the environment.

In [23]:

```
env.close()
```